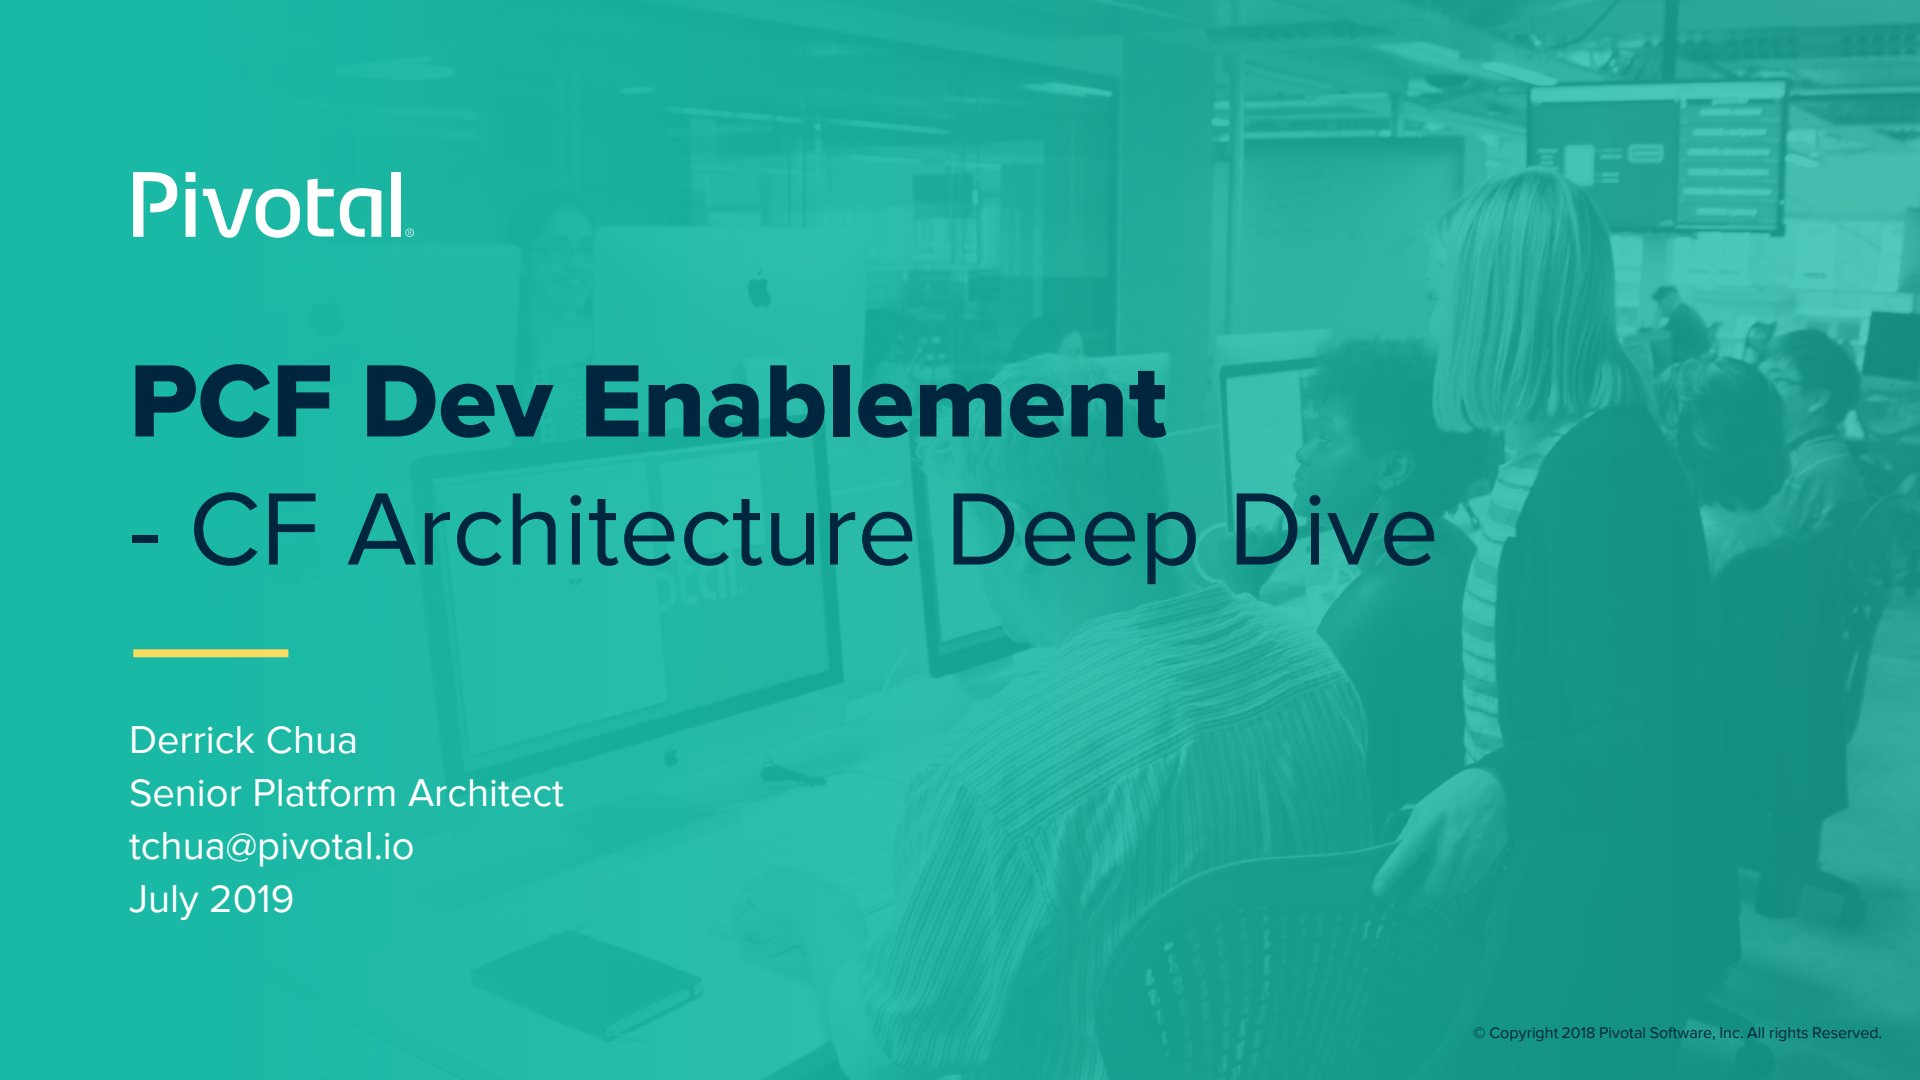

Pivotal

PCF Dev Enablement
- CF Architecture Deep Dive

Derrick Chua
Senior Platform Architect
tchua@pivotal.io
July 2019

© Copyright 2018 Pivotal Software, Inc. All rights Reserved.

PAS vs Opensource CF

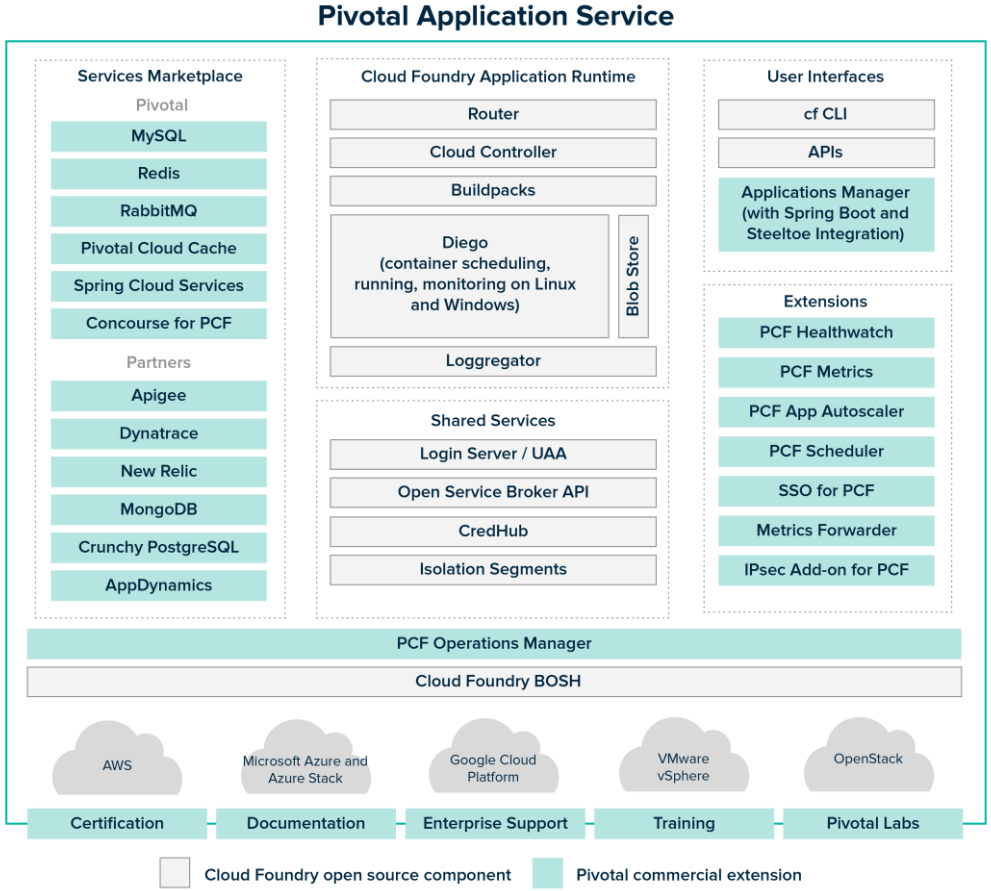

CF Runtime Components

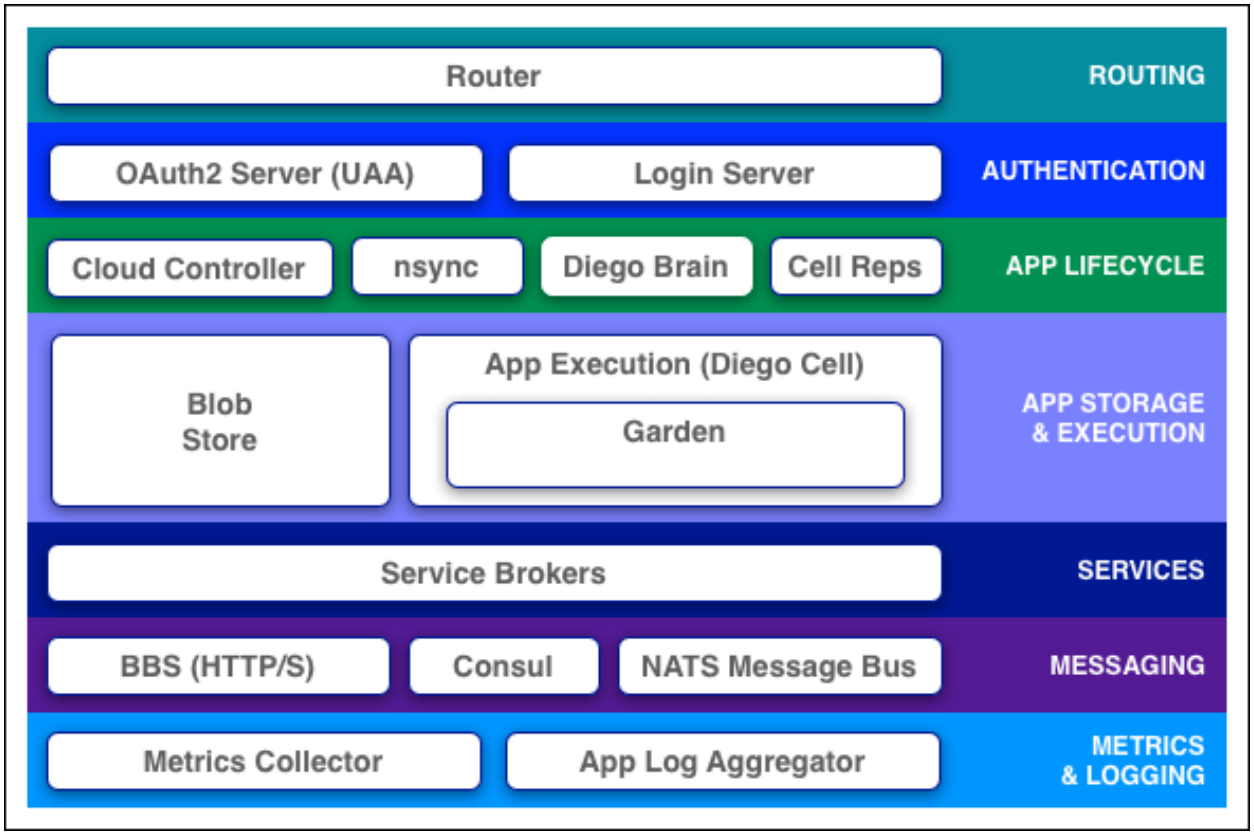

Details at https://docs.pivotal.io/pivotalcf/2-6/concepts/architecture/index.html

nsync, BBS and Cell Reps

Monitor and reconcile application states

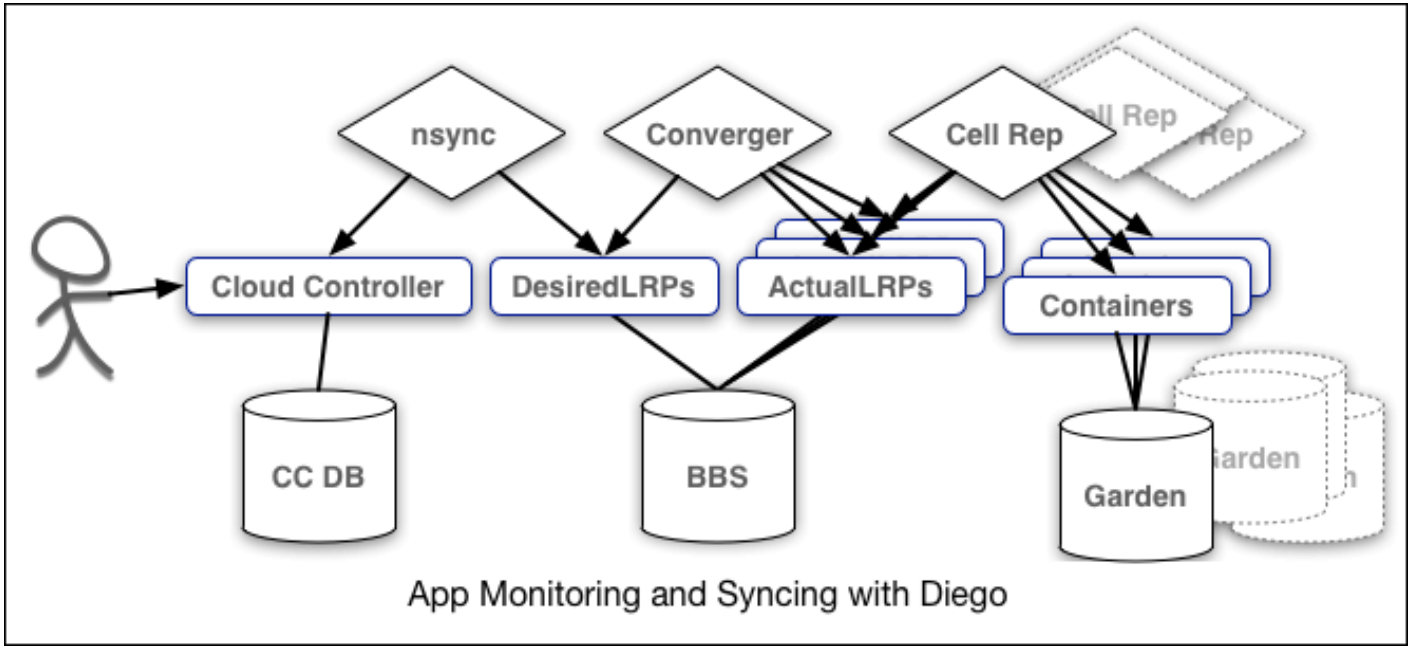

How Diego runs an app

Step 1: Receives request to run an app

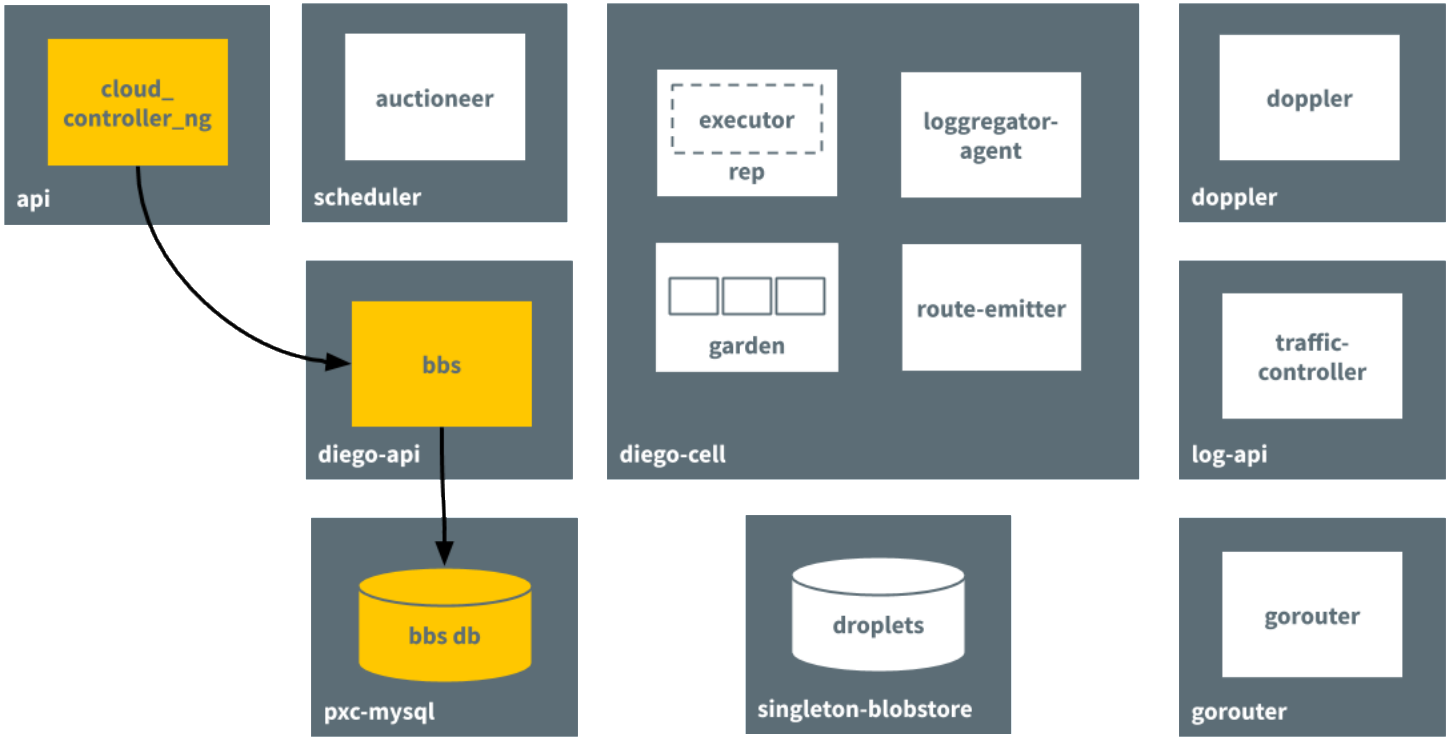

How Diego runs an app

Step 2: Pass request to auctioneer process

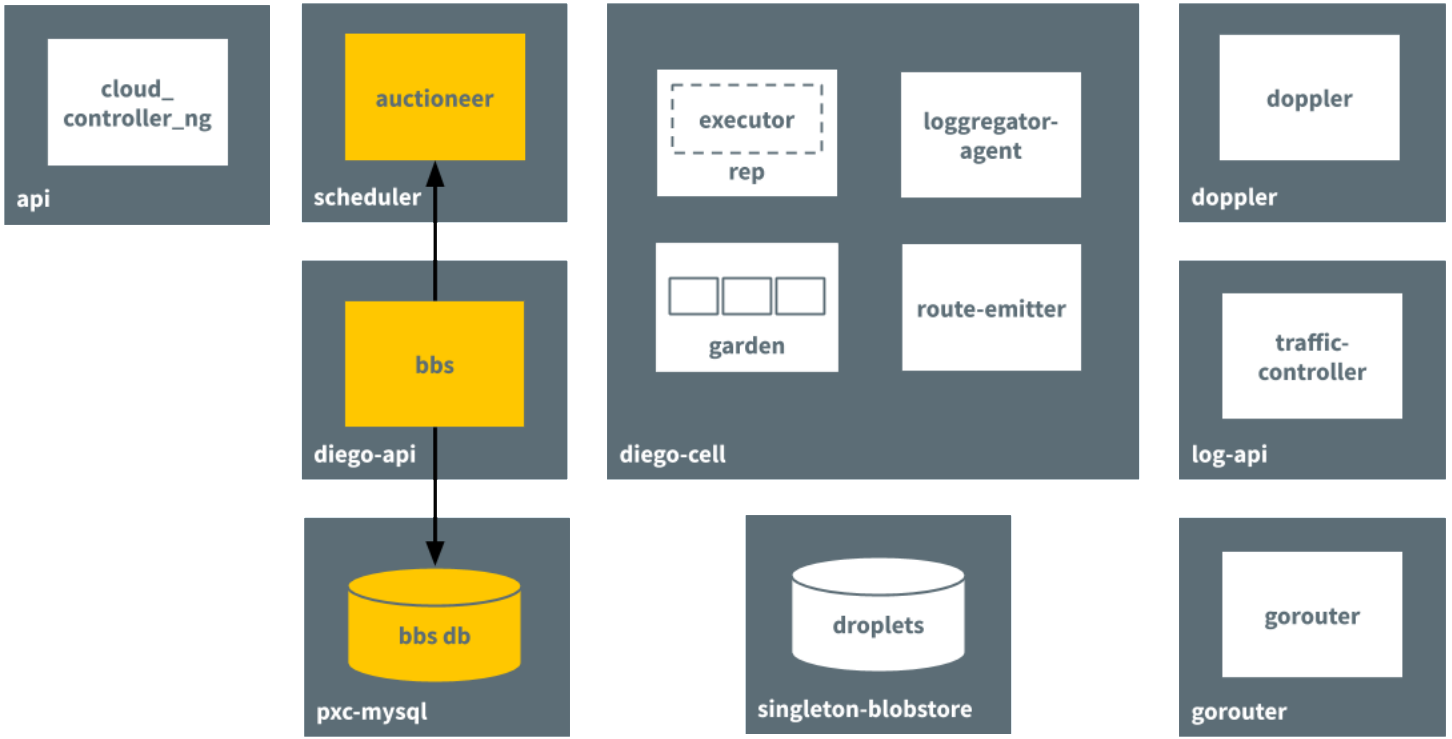

How Diego runs an app

Step 3: Performs auction

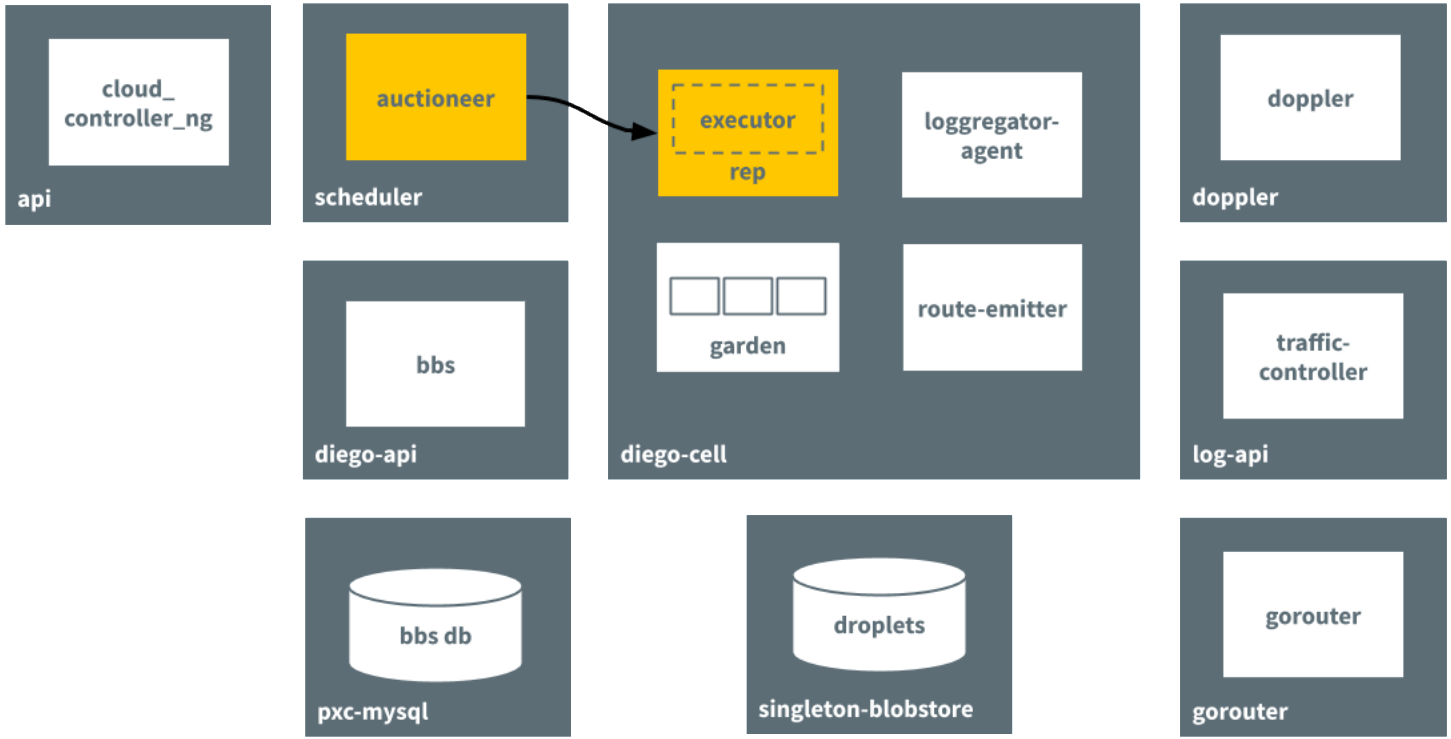

How Diego runs an app

Step 4: Creates container and runs app

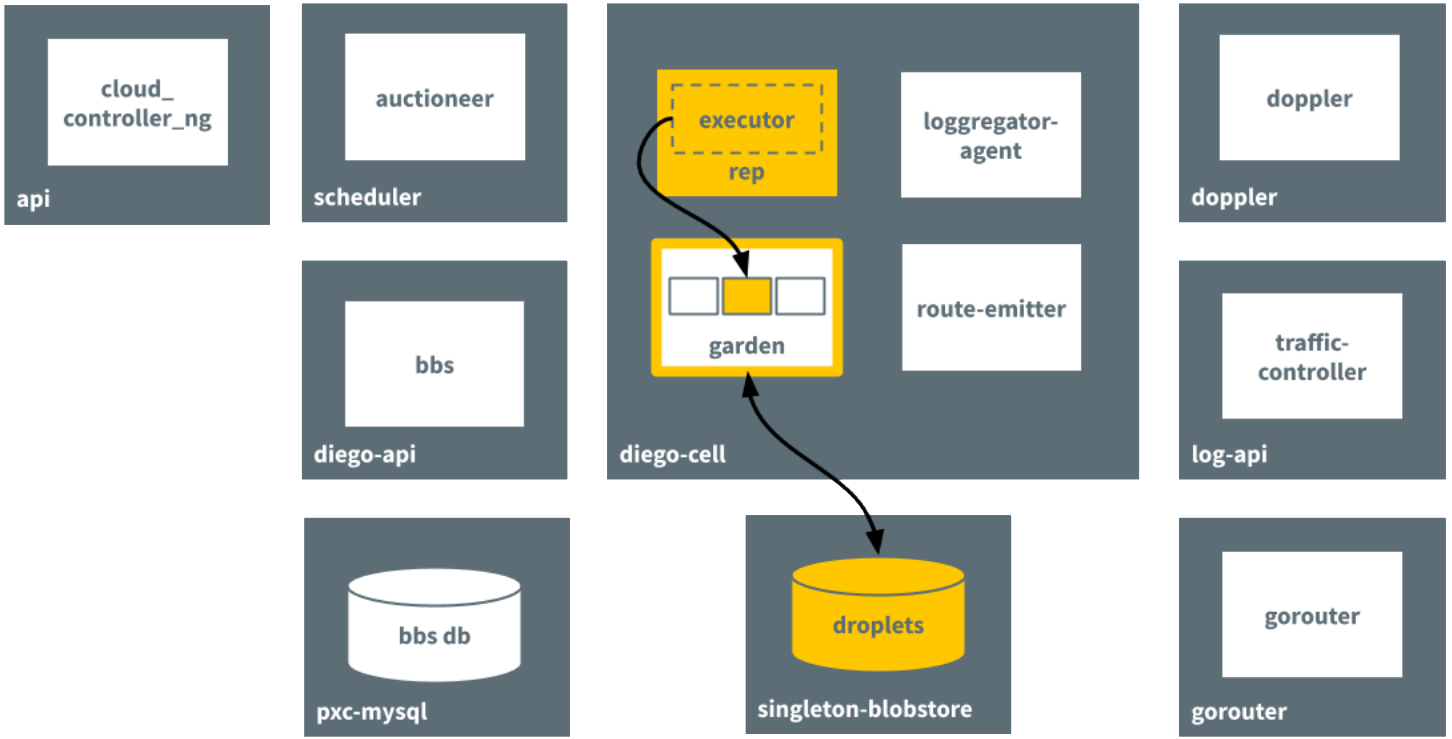

How Diego runs an app

Step 5: Emits routes for app

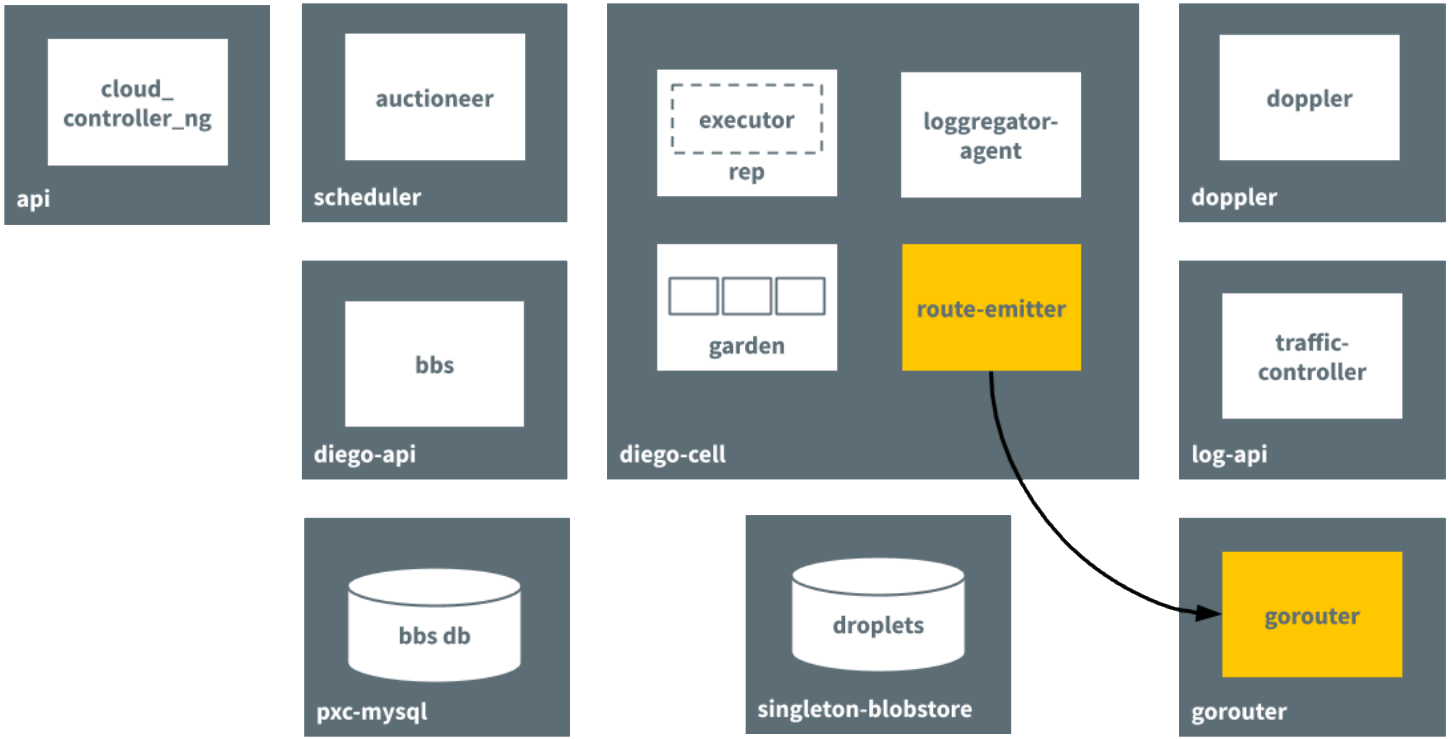

How Diego runs an app

Step 6: Sends logs to loggregator

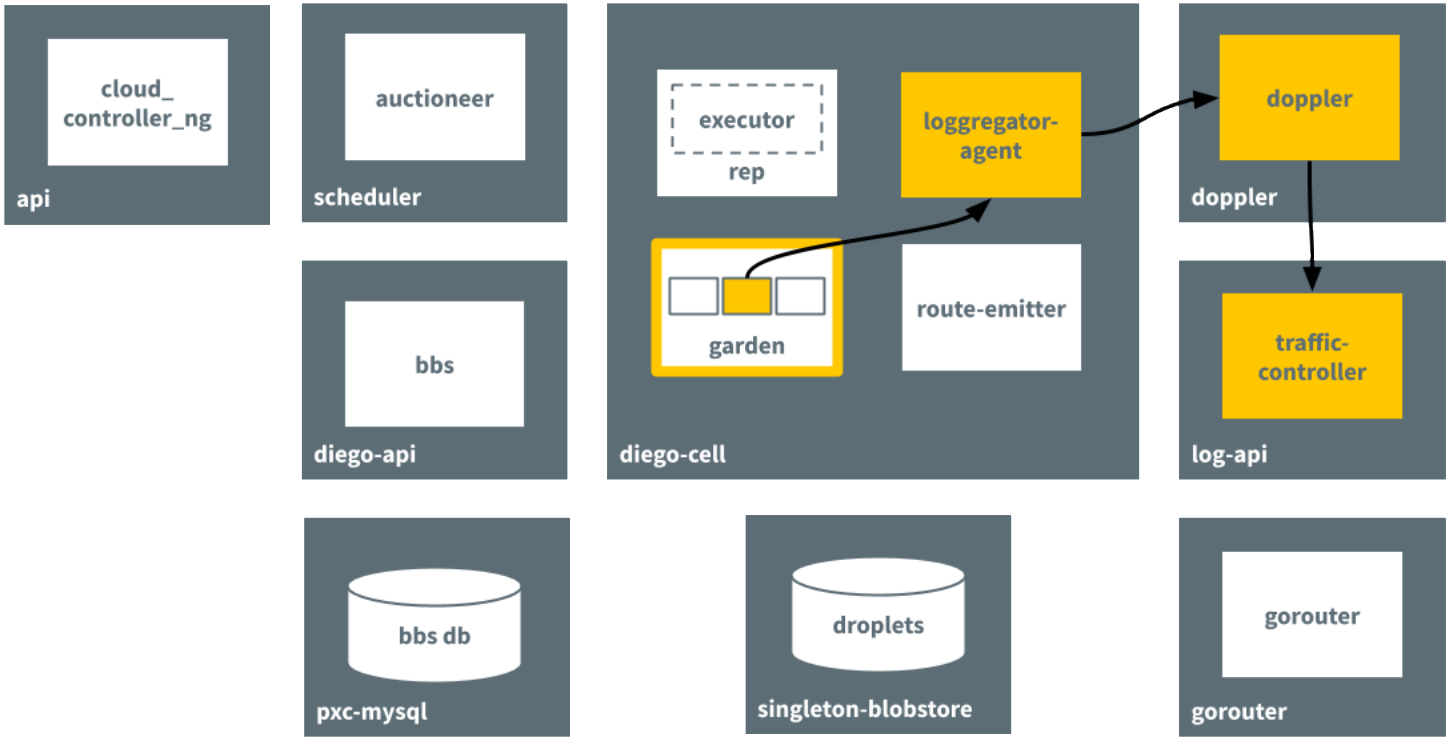

Pivotal

External client request flow

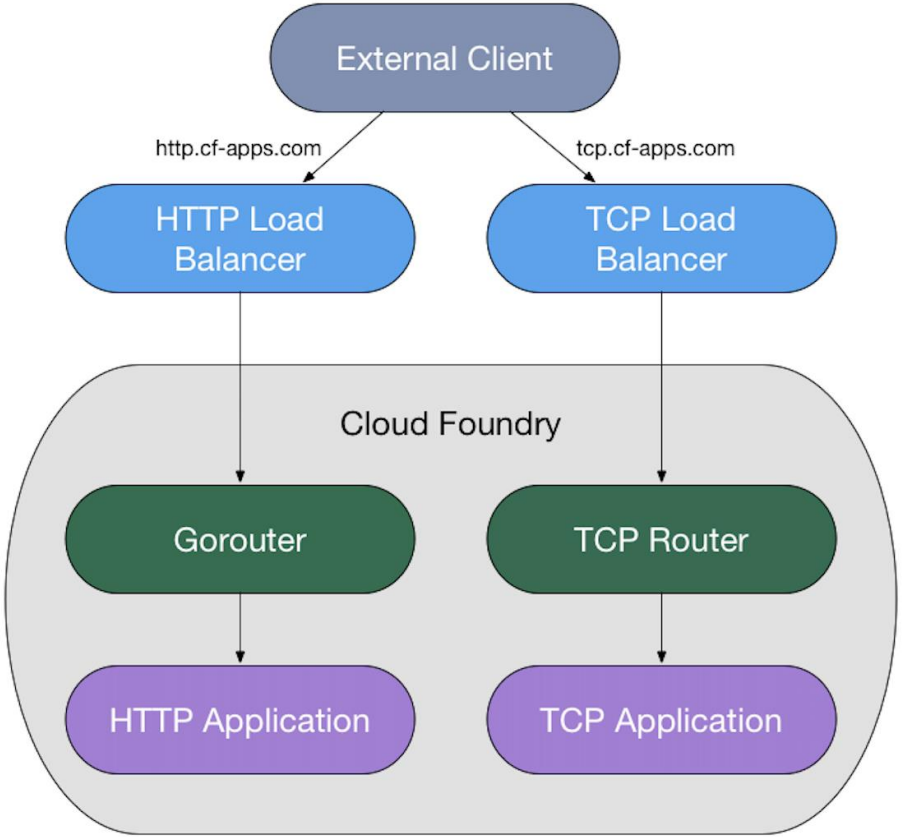

Maintaining updated routing table (HTTP)

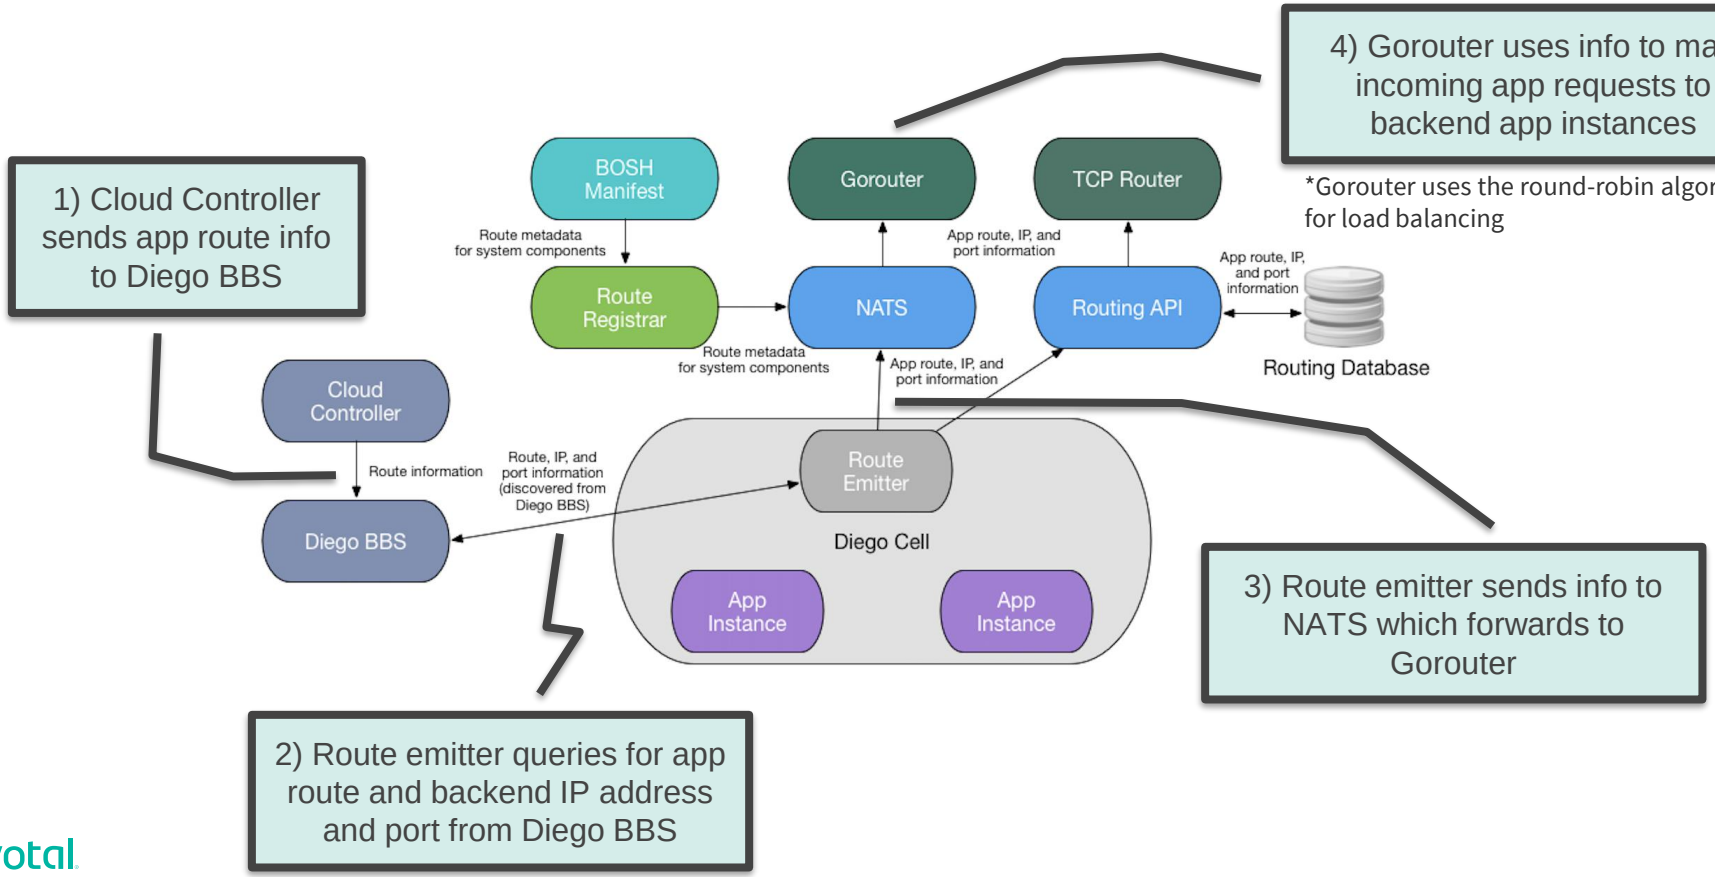

HTTP Routing

Headers

Header name	Purpose
X-Forwarded-Proto	• Gives the scheme of the HTTP request from the client • HTTP for insecure request, HTTPS for secure request • Multiple values – comma separated list • App should process to reject insecure requests
X-Forwarded-For	• Load balancer IP address
X-B3-TraceId X-B3-SpanId	• Zipkin tracing • Logged to Gorouter logs
X-CF-APP-INSTANCE	• Used to obtain debug data for a specific instance of an app
X-Forwarded-Client-Cert	• For mutual TLS • Used to pass the originating client certificate along the data path to the application • If LB terminates TLS, this header should be stripped to prevent client spoofing

** Gorouter has a limit of 1 MB for HTTP Headers

HTTP Routing

Session affinity (sticky sessions)

- To support sticky sessions, configure your app to return a `JSESSIONID` cookie in responses. The app generates a `JSESSIONID` as a long hash in the following format:

```
1A530637289A03B07199A44E8D531427
```

- If an app returns a `JSESSIONID` cookie to a client request, the CF routing tier generates a unique `VCAP_ID` for the app instance based on its GUID in the following format:

```
323f211e-fea3-4161-9bd1-615392327913
```

- On subsequent requests, the client must provide both the `JSESSIONID` and `VCAP_ID` cookies.

The CF routing tier uses the `VCAP_ID` cookie to forward client requests to the same app instance every time. The `JSESSIONID` cookie is forwarded to the app instance to enable session continuity. If the app instance identified by the `VCAP_ID` crashes, the Gorouter attempts to route the request to a different instance of the app. If the Gorouter finds a healthy instance of the app, it initiates a new sticky session.

Container to Container Networking

What happens without container to container networking

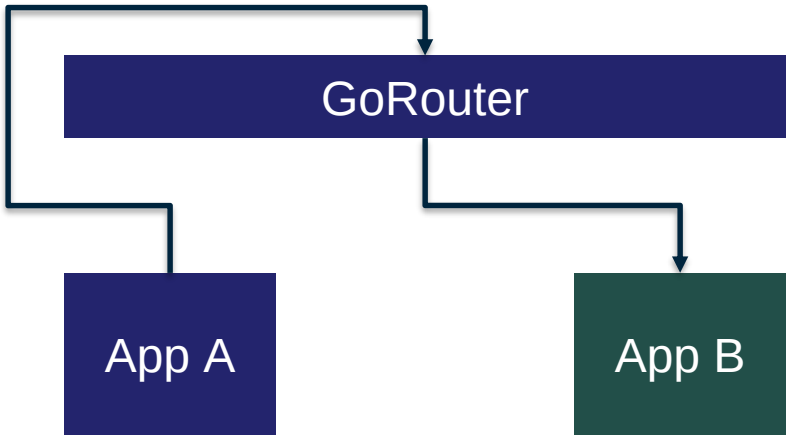

Container to Container Networking

What happens with container to container networking

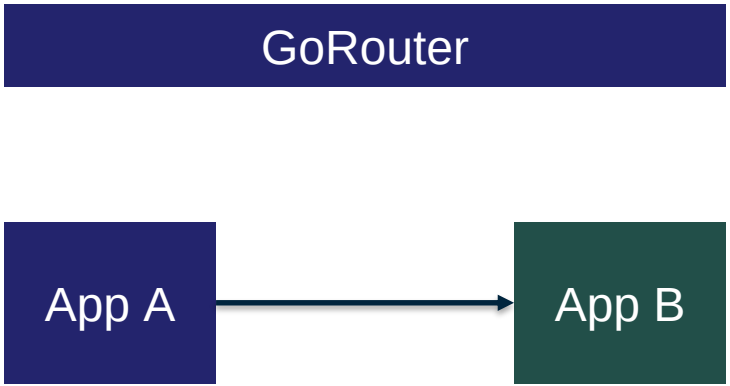

Container to Container Networking

How to...

- By default, each Diego cell in the overlay network is allocated a /24 range that supports 254 containers per cell, one container for each of the usable IP addresses

- Add a network policy to allow one app to talk to another

```
cf add-network-policy SOURCE_APP --destination-app DESTINATION_APP -s
DESTINATION_SPACE_NAME -o DESTINATION_ORG_NAME --protocol (tcp | udp) --port RANGE
```

- With **app service discovery**, apps pushed to Pivotal Application Service (PAS) can establish container-to-container communications through a known route served by internal BOSH DNS
  - Default internal domain is apps.internal
  - PAS apps can reach each other through [APP_NAME].apps.internal

Application Security Groups

- Collection of **egress rules** that specify the protocols, ports, and IP address ranges where app or task instances send traffic

- Define **allow** rules, and their order of evaluation is unimportant when multiple ASGs apply to the same space or deployment

- Administrators can define **a staging ASG** for app and task staging, and a **running ASG** for app and task runtime
    - Staging ASG is typically less restrictive and is used to pull resources required during staging

- Administrators can assign **platform-wide ASGs** that apply to all app and task instances for the entire deployment, or **space-scoped ASGs** that apply only to apps and tasks in a particular space

Application Security Groups

Creating and binding ASGs

```
$ cf create-security-group my-asg ~/workspace/my-asg.json
```

```
[
  {
    "protocol": "icmp",
    "destination": "0.0.0.0/0",
    "type": 0,
    "code": 0
  },
  {
    "protocol": "tcp",
    "destination": "10.0.11.0/24",
    "ports": "80,443",
    "log": true,
    "description": "Allow http and https traffic to ZoneA"
  }
]
```

```
$ cf bind-security-group my-asg my-org my-space
```

Pivotal

Details at https://docs.pivotal.io/pivotalcf/2-6/concepts/asg.html

Container to Container Networking vs ASGs

	ASGs	Container-to-Container Networking Policies
Policy granularity	From a space to an IP address range	From a source app to a destination app
Scope	For a space, org, or deployment	For app to app only
Traffic direction	Outbound control	Policies apply for incoming packets from other app instances
Source app	Is not known	Is identified because of direct addressability
Policies take affect	After app restart	Immediately


The (Near) Future

Transforming How The World Builds Software

© Copyright 2018 Pivotal Software, Inc. All rights Reserved.